

# When Models Meet Data

Liang Zheng


Australian National University

[liang.zheng@anu.edu.au](mailto:liang.zheng@anu.edu.au)

# Spin parity of spiral galaxies II: a catalogue of 80 k spiral galaxies using big data from the Subaru Hyper Suprime-Cam survey and deep learning


(Tosaki et al., Monthly Notices of the Royal Astronomical Society, 2020)

- 560,000 galaxies have been detected in the images captured by the Subaru Telescope
- an accuracy of 97.5%
- it identified spirals in about 80,000 galaxies



# 8.1 Data, Models, and Learning

- A machine learning system has three major components:
- Data, models, learning
- A **model** is obtained by **learning** from the training **data**
- A **prediction** is made by applying a learned **model** on test **data**



# 8.1 Data, Models, and Learning

- We aim to learn **good** models.
- How is **good** defined? We need to have performance metrics on the test data. Examples include
  - Classification accuracy
  - Distance from the ground truth
  - Test time (efficiency)
  - Model size
  - .....
- New performance metrics are constantly being proposed by the machine learning community.

## 8.1.1 Data as Vectors

- Data, read by computers, should be in a numerical format.
- See the tabular format below

Name	Gender	Degree	Postcode	Age	Annual salary
Aditya	M	MSc	W21BG	36	89563
Bob	M	PhD	EC1A1BA	47	123543
Chloé	F	BEcon	SW1A1BH	26	23989
Daisuke	M	BSc	SE207AT	68	138769
Elisabeth	F	MBA	SE10AA	33	113888

- Row: an instance
- Column: a particular feature
- Apart from tabular format, machine learning can be applied to many types of data, e.g., genomic sequences, text and image contents of a webpage, and social media graphs, citation networks...

- We convert the table into numerical format

Name	Gender	Degree	Postcode	Age	Annual salary
Aditya	M	MSc	W21BG	36	89563
Bob	M	PhD	EC1A1BA	47	123543
Chloé	F	BEcon	SW1A1BH	26	23989
Daisuke	M	BSc	SE207AT	68	138769
Elisabeth	F	MBA	SE10AA	33	113888



Gender ID	Degree	Latitude (in degrees)	Longitude (in degrees)	Age	Annual Salary (in thousands)
-1	2	51.5073	0.1290	36	89.563
-1	3	51.5074	0.1275	47	123.543
+1	1	51.5071	0.1278	26	23.989
-1	1	51.5075	0.1281	68	138.769
+1	2	51.5074	0.1278	33	113.888

- Gender is quantized to -1 and +1
- Degree from BS, MS to PhD: 1, 2, 3
- Postcode corresponds to Latitude and Longitude on the map
- Name is removed because of privacy and because it does not contain useful information for the machine learning system. (exceptions? See [1])

[1] Chen et al., What's in a Name? First Names as Facial Attributes. CVPR 2013

- We use  $N$  to denote the number of examples in a dataset and index the examples with lowercase  $n = 1, \dots, N$

Gender ID	Degree	Latitude (in degrees)	Longitude (in degrees)	Age	Annual Salary (in thousands)
-1	2	51.5073	0.1290	36	89.563
-1	3	51.5074	0.1275	47	123.543
+1	1	51.5071	0.1278	26	23.989
-1	1	51.5075	0.1281	68	138.769
+1	2	51.5074	0.1278	33	113.888

- Each row is a particular individual  $x_n$  referred to as an **example** or **data point** in machine learning
- The subscript  $n$  refers to the fact that this is the  $n$ th example out of a total of  $N$  examples in the dataset
- Each column represents a particular feature of interest about the example, and we index the features as  $d = 1, \dots, D$
- Each example is a  $D$ -dimensional vector


- Consider the problem of predicting annual salary from age

$D$  columns

	Gender ID	Degree	Latitude (in degrees)	Longitude (in degrees)	Age	Annual Salary (in thousands)
$N$ rows	-1	2	51.5073	0.1290	36	89.563
	-1	3	51.5074	0.1275	47	123.543
	+1	1	51.5071	0.1278	26	23.989
	-1	1	51.5075	0.1281	68	138.769
	+1	2	51.5074	0.1278	33	113.888

- A supervised learning algorithm
- We have a label  $y_n$  (the salary) associated with each example  $x_n$  (age).
- A dataset is written as a set of example-label pairs  $\{(x_1, y_1), \dots, (x_n, y_n), \dots, (x_N, y_N)\}$
- The table of examples  $\{x_1, \dots, x_N\}$  are concatenated and written as  $X \in \mathbb{R}^{N \times D}$

We are interested in:  
 What is the salary ( $y$ ) at  
 age 60 ( $x = 60$ )?





## 8.1.2 Models as Functions

- Once we have data in an appropriate vector representation, we can construct a **predictive function** (known as a **predictor**).
- Here, a model means a **predictor**.
- A predictor is a function that, when given a particular input example (in our case, a vector of features), produces an output.
- For example,

$$f: \mathbb{R}^D \rightarrow \mathbb{R}$$


where the input  $\mathbf{x}$  is a  $D$ -dimensional vector, and the output is a real-valued scalar. That is, the function  $f$  is applied to  $\mathbf{x}$ , written as  $f(\mathbf{x})$  and returns a real number.

## 8.1.2 Models as Functions

- We mainly consider the special case of linear functions

$$f(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x} + \theta_0$$

- Example: predicting salary  $f(\mathbf{x})$  from age  $\mathbf{x}$ .




Black and solid diagonal line is an example predictor.

$$f(60) = 100$$

## 8.1.3 Models as Probability Distributions

- The observed data is usually a combination of the **true underlying data** and **noise**, i.e.,  $\tilde{x} = x + n$
- We wish to reveal  $x$  from  $\tilde{x}$
- So We would like to have predictors that express some sort of uncertainty, e.g., to quantify the confidence we have about the value of the prediction for a particular test data point.



Example function (black solid diagonal line) and its predictive uncertainty at  $x = 60$  (drawn as a Gaussian).

- Instead of considering a predictor as a single function, we could consider predictors to be **probabilistic models**.
- We will learn probability in later lectures

# 8.1.4 Learning is Finding Parameters

- The goal of learning is to find a **model** and its corresponding **parameters** such that the resulting **predictor** will perform well on unseen **data**.
- 3 algorithmic phases when discussing machine learning algorithms
- Prediction or inference
- Training or parameter estimation
- Hyperparameter tuning or model selection
- **Prediction phase**: we use a trained predictor on previously unseen test data
- **The training or parameter estimation phase**: we adjust our predictive model based on training data. We will introduce the **empirical risk minimization** for finding good parameters.
- We use **cross-validation** to assess predictor performance on unseen data.
- We also need to balance between fitting well on training data and finding “simple” explanations of the phenomenon. This trade-off is often achieved using **regularization**.

- Hyperparameter tuning or model selection
  - We need to make high-level modeling decisions about the structure of the predictor. For example
  - Number of layers to be used in deep learning
  - Number of components in a Gaussian Mixture Model
  - Weight of regularization terms
- } Hyperparameter
- The problem of choosing among different models/hyperparameters is called **model selection**
  - Difference between **parameters** and **hyperparameters**
  - Parameters are to be numerically optimized ( $\sim 10^6$  weights in a deep network)
  - Hyperparameters need to use search techniques (neural architecture search [2])

## 8.2 Empirical Risk Minimization

- What does it mean to **learn**?
- Estimating parameters based on training data.
- Four questions will be answered
- What is the set of functions we allow the predictor to take? – **Hypothesis class of functions**
- How do we measure how well the predictor performs on the training data? -- **Loss functions for training**
- How do we construct predictors from only training data that performs well on unseen test data? -- **regularization**
- What is the procedure for searching over the space of models? - **Cross-Validation**

## 8.2.1 Hypothesis Class of Functions

- We are given  $N$  examples  $\mathbf{x}_n \in \mathbb{R}^D$  and corresponding scalar labels  $y_n \in \mathbb{R}$ .
- Supervised learning: we have pairs  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$
- We want to estimate a predictor  $f(\cdot, \boldsymbol{\theta}): \mathbb{R}^D \rightarrow \mathbb{R}$ , parametrized by  $\boldsymbol{\theta}$
- We hope to be able to find a good parameter  $\boldsymbol{\theta}^*$  such that we fit the data well, that is

$$f(\mathbf{x}_n, \boldsymbol{\theta}^*) \approx y_n \text{ for all } n = 1, \dots, N$$

- We use  $\hat{y}_n = f(\mathbf{x}_n, \boldsymbol{\theta}^*)$  to represent the output of the predictor

# Example (least-squares regression)

- When the label  $y_n$  is real-valued, a popular choice of function class for predictors is **affine functions** (linear functions).

$$f(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x} + \theta_0$$

- For more compact representations, we concatenate an additional unit feature  $x^{(0)} = 1$  to  $\mathbf{x}_n$ , i.e.,

$$\mathbf{x}_n = [x_n^{(0)}, x_n^{(1)}, x_n^{(2)}, \dots, x_n^{(D)}]^T = [1, x_n^{(1)}, x_n^{(2)}, \dots, x_n^{(D)}]^T$$

- The parameter vector is  $\boldsymbol{\theta} = [\theta_0, \theta_1, \theta_2, \dots, \theta_D]^T$
- We can write the predictor as follows

$$f(\mathbf{x}_n, \boldsymbol{\theta}) = \boldsymbol{\theta}^T \mathbf{x}_n$$

which is equivalent to the affine model

$$f(\mathbf{x}_n, \boldsymbol{\theta}) = \theta_0 + \sum_{d=1}^D \theta_d x_n^{(d)} = \theta_0 x_n^{(0)} + \sum_{d=1}^D \theta_d x_n^{(d)} = \boldsymbol{\theta}^T \mathbf{x}_n$$



# Example (least-squares regression)

$$f(\mathbf{x}_n, \boldsymbol{\theta}) = \boldsymbol{\theta}^T \mathbf{x}_n$$

- The predictor takes the vector of features representing a single example  $\mathbf{x}_n$  as input and produces a real-valued output,

$$f: \mathbb{R}^{D+1} \rightarrow \mathbb{R}$$

- $f(\mathbf{x}_n, \boldsymbol{\theta}) = \boldsymbol{\theta}^T \mathbf{x}_n$  is a linear predictor
- There are many non-linear predictors, such as the neural networks

## 8.2.2 Loss Function for Training

- In training, we aim to learn a model that **fits the data well**.
- To define “**fits the data well**”, we specify a **loss function**  
 $\ell(y_n, \hat{y}_n)$
- Input: ground truth label  $y_n$  of a training example  
the prediction  $\hat{y}_n$  of this training example
- Output: a non-negative number, called **loss**. It represents how much error we have made on this particular prediction
- To find good parameters  $\theta^*$ , we need to minimize the average loss on the set of  $N$  training examples
- We usually assume training examples  $(x_1, y_1), \dots, (x_N, y_N)$  are **independent and identically distributed (i.i.d)**.

- Under the i.i.d assumption, the empirical mean is a good estimate of the population mean.
- We can use the empirical mean of the loss on the training data
- Given a **training set**  $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$ , we use the notation of an example matrix

$$\mathbf{X} := [\mathbf{x}_1, \dots, \mathbf{x}_N]^T \in \mathbb{R}^{N \times D}$$

and a label vector

$$\mathbf{y} = [y_1, \dots, y_N]^T \in \mathbb{R}^N$$

- The average loss is given by

$$\mathbf{R}_{emp}(f, \mathbf{X}, \mathbf{y}) = \frac{1}{N} \sum_{n=1}^N \ell(y_n, \hat{y}_n)$$

where  $\hat{y}_n = f(\mathbf{x}_n, \boldsymbol{\theta})$ . The above equation is called the **empirical risk**. The learning strategy is called **empirical risk minimization**.

# Example - Least-Squares Loss

- We use the squared loss function

$$\ell(y_n, \hat{y}_n) = (y_n - \hat{y}_n)^2$$

- We aim to minimize the empirical risk, which is the average of the losses over the training data.

$$\min_{\boldsymbol{\theta} \in \mathbb{R}^D} \frac{1}{N} \sum_{n=1}^N \ell(y_n, \hat{y}_n) = \min_{\boldsymbol{\theta} \in \mathbb{R}^D} \frac{1}{N} \sum_{n=1}^N (y_n - f(\mathbf{x}_n, \boldsymbol{\theta}))^2$$

Using the linear predictor  $f(\mathbf{x}_n, \boldsymbol{\theta}) = \boldsymbol{\theta}^T \mathbf{x}_n$ , we obtain the optimization problem

$$\min_{\boldsymbol{\theta} \in \mathbb{R}^D} \frac{1}{N} \sum_{n=1}^N (y_n - f(\mathbf{x}_n, \boldsymbol{\theta}))^2$$

- This equation can be equivalently expressed in matrix form

$$\min_{\boldsymbol{\theta} \in \mathbb{R}^D} \frac{1}{N} \|\mathbf{y} - \mathbf{X}\boldsymbol{\theta}\|^2$$

- This is known as the **least-squares problem**. There exists a closed-form analytic solution for this by solving the normal equations. We will discuss it in later lectures

- We actually want to find a predictor  $f$  that minimizes the **expected risk** (or the population risk)

$$\mathbf{R}_{\text{true}}(f) = \mathbb{E}_{x,y}[\ell(y, f(x))]$$

where  $y$  is the ground truth label and  $f(x)$  is the prediction based on the example  $x$ .

- $\mathbf{R}_{\text{true}}(f)$  is the true risk, if we can access an infinite amount of data
- The expectation  $\mathbb{E}$  is over the infinite set of all possible data and labels.

- Machine learning applications have different types of performance measure.
  - For classification: accuracy, AUC, F1 score, etc.
  - For detection: mean average precision, mIoU, etc.
  - For image denoise/super resolution: SSIM, PSNR, etc.
- In principle, the loss function should correspond to the measure.
- However, there are often mismatches between loss functions and the measures – due to implementation/optimization considerations

# Check your understanding

T

- A machine learning model may contain as few as a couple of parameters

T

- When we use a linear regression modeling,  $f(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x} + \theta_0$ , we don't have hyperparameters.

F

- Hyperparameters are usually learned through the same way as normal parameters.

T

- It's very hard to know the expected risk, but easier to know the empirical risk

- Given a fixed task, we can only use a fixed set of evaluation

F

metrics.